

## APAI2023 - LAB03

# PULP\_Embedded\_Programming

*Authors: Lorenzo Lamberti, Manuele Rusci, Francesco Conti*

*Contacts: [lorenzo.lamberti@unibo.it](mailto:lorenzo.lamberti@unibo.it), [d.nadalini@unibo.it](mailto:d.nadalini@unibo.it), [alberto.dequino@unibo.it](mailto:alberto.dequino@unibo.it)*

**Links:** [GitHub Link \(code\)](#) [GDOC link \(assignment\)](#)

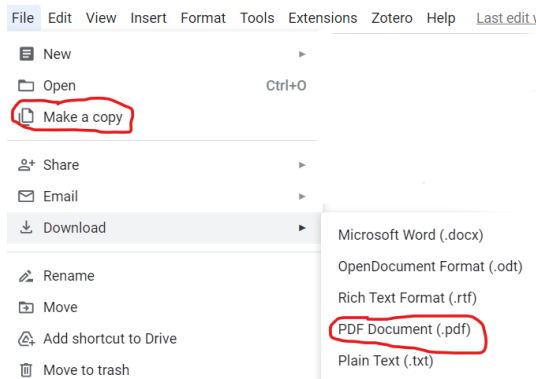
## Summary

1. Subject(s):
  - PULP architecture, vector sum, matrix-vector mul, profiling code execution;
2. Programming Language: C;
3. Lab duration: 3h
4. Objective: Embedded programming & profiling  
you will learn basics of embedded programming, the pulp architecture, basic operations (sum & matmul), and how to profile your code execution (MAC, cycles) !
5. Assignment:
  - 2 tasks
  - Time for delivery: 1 week
  - **Submission deadline: Oct 26<sup>th</sup> 2023**

## How to deliver the assignment

You will deliver ONLY THIS TEXT FILE, no code

- Copy this google doc to your drive, so that you can modify it. (File -> make a copy)
- Fill the tasks on this google doc.
- Export to pdf format.
- Rename the file to: LAB<number\_of\_the\_lesson>\_APAI\_<your\_name>.pdf
- Use Virtuale platform to load ONLY your .pdf file



## LAB STARTS HERE

### 0. Access to the remote server, and setup

- Open this web page: <https://compute.eees.dei.unibo.it:8443/guacamole/>  
**(works only from ALMA WIFI NETWORK!)**
- Login. We distribute credentials by hand.
- Open a terminal (right click – open a new terminal)
- Clone: `git clone https://github.com/EEESlab/APAI23-LAB03-PULP-Embedded-Programming`
- `module load pulp-sdk`
- `cd APAI23-LAB03-PULP-Embedded-Programming`
- `cd pulp-helloworld`
- `make clean all run`

### Task 1: vector sum

#### 0. Setup:

- Open VSCode.

- Go to “vector\_sum/” folder
- Every time you want to run the code, **SAVE your file** and write in the terminal “make clean all run”

### 1.1. Define N to 50 (N= vector size) and run the code:

	Question	Answer
1	What's the result of the vector_sum() function?	
2	Is the checksum correct ?	
3	Print all the elements of “array_1”. What's the output?	

#### **Tips:**

- **Question 3:** to print the element of the array, uncomment the function call “print\_array(array\_1, N)” inside the main()

### 1.2. Define N to 350 (N= vector size) and run the code:

	Question	Answer
1	What's the result of the vector_sum() function?	
2	Is the checksum correct ?	
3	Print all the elements of “array_1”. What's the output?	

4	Array_1 should be filled with increasing values from 1 to N. Why isn't the case here?	
---	---	--

**Tips:**

- **Question 3:** to print the element of the array, uncomment the function call "print\_array(array\_1, N)" inside the main()
- **Question4:** what range of values a "char" data type can represent?

**1.3. Fix the issue: cast "array\_1" and the function's arguments to int (or short int)**

Write down your solution


before	After your fix
<code>unsigned char array_1[N];</code>	
<code>int init_array(unsigned char * A_ar, int size)</code>	
<code>void print_array(unsigned char * A_ar, int size)</code>	
<code>int vector_sum(unsigned char * A_ar, int size)</code>	

After the fix, answer again

	Question	Answer
1	What's the result of the vector_sum() function?	
2	Is the checksum correct ?	
3	Print all the elements of "array_1".	

What's the output?

## Task 2: matrix-vector product & profiling



### 0. Setup:

- Open VSCode.
- Go to “matrix\_vector\_product/” folder
- Every time you want to run the code, **SAVE your file** and write in the terminal “make clean all run”

### 2.1. Implement missing code

GOAL: We want to count how many Mutiply-ACcumulate operations we perform for a GEMV

TASK: Increment a counter every time we perform a MAC inside the GEMV inner loop. I prepared a variable dallec mac\_counter. This is the reference code

```

int mac_counter=0;
/* generic matrix-vector multiplication */
int __attribute__((noinline)) gemv(int size_N, int size_M, int *mat_i, int *vec_i, int *vec_o)
{
    for (int i = 0; i < size_N; i++) // outer loop of the gemv
    {
        for (int j = 0; j < size_M; j++) // inner loop of the gemv
        {
            // multiply accumulate operation (MAC)
            vec_o[i] += mat_i[i * size_M + j] * vec_i[j];
            /*(vec_o+i) += *(mat_i+i*M+j)**(vec_i+j); // try to uncomment this and comment the above line
        }
    }
}

```

Put your solution below (code)

**[HERE]**

## 2.2. Implement missing code part II

Add performance counters to profile the gemv. The gemv is calculated with this line of code in the main():

```
gemv(N, M, matrix, vector, output_vec);
```

To profile it, exploit these two functions I prepared for you. You should start the profiling right before and stop it right after

```
start_perf_counter();
stop_perf_counter();
```

Put your solution below (code)

**[HERE]**

## 2.3. Implement missing code part III

Enable the performance counters of our interest. We want to profile:

- Execution cycles (total)
- N° instructions executed

Here's the full list of the performance counters

```

typedef enum {
    PI_PERF_CYCLES      = 17, /*!< Total number of cycles (also includes the
        cycles where the core is sleeping). Be careful that this event is using a
        timer shared within the cluster, so resetting, starting or stopping it on
        one core will impact other cores of the same cluster. */
    PI_PERF_ACTIVE_CYCLES = 0, /*!< Counts the number of cycles the core was
        active (not sleeping). */
    PI_PERF_INSTR       = 1, /*!< Counts the number of instructions executed.
        */
    PI_PERF_LD_STALL    = 2, /*!< Number of load data hazards. */
    PI_PERF_JR_STALL    = 3, /*!< Number of jump register data hazards. */
    PI_PERF_IMISS       = 4, /*!< Cycles waiting for instruction fetches, i.e.
        number of instructions wasted due to non-ideal caching. */
    PI_PERF_LD          = 5, /*!< Number of data memory loads executed.
        Misaligned accesses are counted twice. */
    PI_PERF_ST          = 6, /*!< Number of data memory stores executed.
        Misaligned accesses are counted twice. */
    PI_PERF_JUMP        = 7, /*!< Number of unconditional jumps (j, jal, jr,
        jalr). */
    PI_PERF_BRANCH      = 8, /*!< Number of branches. Counts both taken and
        not taken branches. */
    PI_PERF_BTAKEN      = 9, /*!< Number of taken branches. */
    PI_PERF_RVC         = 10, /*!< Number of compressed instructions
        executed. */
    PI_PERF_LD_EXT      = 12, /*!< Number of memory loads to EXT executed.
        Misaligned accesses are counted twice. Every non-TCDM access is considered
        external (cluster only). */
    PI_PERF_ST_EXT      = 13, /*!< Number of memory stores to EXT executed.
        Misaligned accesses are counted twice. Every non-TCDM access is considered
        external (cluster only). */
    PI_PERF_LD_EXT_CYC  = 14, /*!< Cycles used for memory loads to EXT.
        Every non-TCDM access is considered external (cluster only). */
    PI_PERF_ST_EXT_CYC  = 15, /*!< Cycles used for memory stores to EXT.
        Every non-TCDM access is considered external (cluster only). */
    PI_PERF_TCDM_CONT   = 16, /*!< Cycles wasted due to TCDM/log-interconnect
        contention (cluster only). */
} pi_perf_event_e;

```

Ref: `/rtos/pmsis/pmsis_api/include/pmsis/chips/default.h`

You can enable them in the `start_perf_counter()` function.

Complete the code where you find `/* YOUR CODE HERE */` with the right performance counters (see previous figure)

```

void start_perf_counter()
{
    // enable the perf counter of interest
    pi_perf_conf( 1 << PI_PERF_/* YOUR CODE HERE*/ | // count cycles
                1 << PI_PERF_/* YOUR CODE HERE*/ ); // count instructions
    // reset the perf counters
    pi_perf_reset();
    // start the perf counter
    pi_perf_start();
}

```

Put your solution below (code)

**[HERE]**

## 2.4. Implement missing code part IV

GOAL: we now are able to measure Cycles & instructions (thanks to performance counters), and MAC operations (your implementation in the code)

TASK: calculate and print the following metrics in the code (fill in where you read `/*YOUR CODE HERE*/`)

- CPI = cycles/instructions
- MAC/cycles
- Instructions/Cycles
- Instructions/MACs

```
/*
TASK 2.2
Measure:
- How many MAC operations are needed for the gemv
Calculate:
- CPI
- MACs/Cycles
- Instructions/Cycles
- Instructions/ MACs
*/

// N° Multiply Accumulate Operations (MACs)
/* already done with task 2.1. use the mac counter global variable */
// CPI = cycles / n°instructions_executed
float cpi = /* YOUR CODE HERE*/
// MAC/Cycles
float mac_on_cycles = /* YOUR CODE HERE*/
// Instructions/Cycles
float instructions_on_cycles = /* YOUR CODE HERE*/
// Instructions/MAC
float instructions_on_mac = /* YOUR CODE HERE*/
```

Put your solution below (code)


**[HERE]**

**2. The size of the matrix is NxM=50x50 and the vector is M=50. Run the code and fill the table:**

Note:

- Profile the matrix-vector product with different compiler optimizations: -O1, -O3, -O3 with HW Loops (default is -O1)





	<b>-01</b>	<b>-03 -mnohwloops</b>	<b>-03</b>
Cycles			
N°Instructions			
MACs			
CPI			
Instructions/Cycles			
Instructions/MACs			

